

Eliciting Trust Values from Recommendation Errors

John O'Donovan, Barry Smyth

Department of Computer Science, University College Dublin, Ireland
john.odonovan@ucd.ie, barry.smyth@ucd.ie

ABSTRACT

Increasing availability of information has furthered the need for recommender systems across a variety of domains. These systems are designed to tailor each user's information space to suit their particular information needs. Collaborative filtering is a successful and popular technique for producing recommendations based on similarities in users' tastes and opinions. Our work focusses on these similarities and the fact that current techniques for defining which users contribute to recommendation are in need of improvement.

In this paper we propose the use of trustworthiness as an improvement to this situation. In particular, we define and empirically test a technique for eliciting trust values for each producer of a recommendation based on that user's history of contributions to recommendations.

We present three computational models for leveraging under/overestimate errors in users' past contributions to recommendations to generate a recommendation *range* each side of a fixed point on the recommendation scale to be presented to the target user. We show how this trust-based technique can be easily incorporated into a standard collaborative filtering algorithm and define a fair comparison in which our technique outperforms a benchmark algorithm in predictive accuracy.

Keywords

Recommender systems, collaborative filtering, profile similarity, trust

INTRODUCTION

A large amount of research effort has been directed at the well-known problem of information overload, whereby the task of locating relevant information quickly is becoming ever more arduous with increasing volumes of online information (Sinha & Swearingen 2002), (Resnick *et al.* 1994). This is the domain of Recommender Systems, which are deployed in a variety of applications to assist the user in locating the right information at the right time. These systems can be found in diverse areas such as virtual shopping assistants (Dellarocas), restaurant or movie recommenders (Burke, Hammond, & Young 1997) etc. Recommenders

have proven themselves successful at harnessing information normally reserved for social interaction, and leveraging this to provide tailor made solutions for individual users based on the tastes and opinions of similar users. To achieve this solution, recommenders employ ideas from research areas such as user-profiling, information filtering and machine learning.

The majority of research into recommendation strategies focusses on two main approaches. Content-based recommenders operate by comparison of descriptions of content of recommendable items (eg, products or services). A content-based movie recommender would require that a movie had an adequate textual description before it could be incorporated into the system. It would rely on information such as genre, actors, director etc. and match this against the learned preferences of the user to arrive at a suitable recommendation. These systems rely heavily on knowledge engineering, are not suited to domains without a high quality textual description, and do not facilitate serendipitous recommendations, in that the system will always try to recommend items that are similar to those seen and liked before. This is not an accurate reflection of the changing tastes of real world users.

The alternative approach to the recommendation problem is a collaborative one. The collaborative filtering (CF) approach attempts to model the real-world scenario of looking to friends for recommendations (Resnick *et al.* 1994). In this way, the need for specific item-knowledge is eliminated, giving this approach a much more broad and flexible application base. CF systems rely on the availability of rich user profiles containing lots of rating data, so that sufficient profile overlap can be attained to build up user peer groups. Peers, (recommendation partners) for a target user are generated by choosing profiles with similar or have highly correlated ratings histories with that user.

This work focusses on the CF approach to the recommendation task, in ways to improve the accuracy of predictions made by CF approaches, and in ways to increase user-satisfaction and trust in these predictions. We propose to modify the selection process for recommendation partners. Currently this selection is usually based on a profile-profile similarity metric, calculated using some similarity function such as Pearson's correlation coefficient (Resnick *et al.* 1994). We introduce a second metric of *trustworthiness* upon which to base this selection decision. For exam-

ple: People look to their friends for recommendations about items, but some of these friends might not know a lot about that particular item, or may have an esoteric or uncharacteristic opinion about that item, which leads him/her to give consistently skewed or bad recommendations about that particular item. Our proposal is that a recommendation partner should be both similar to the target user, and trustworthy in that the contribution that they have made to previous recommendations has been a positive one. We can model this trustworthiness both at the profile level and at the item level. The former allows us to say how trustworthy a recommendation producer is on a general level, while the latter allows us to provide a specific trust score for a producer with respect to the particular item that user is involved in recommending. For our evaluations in this paper, we use recommendation error history to elicit these trust values. We describe three methods of incorporating these values into a regular CF algorithm, and show results of an empirical comparison of these techniques. We also show that our trust-based recommendation strategy outperforms a standard benchmark system which uses Resnick's prediction algorithm (Resnick *et al.* 1994) over a standard CF dataset.

BACKGROUND

Trust, reputation and reliability are factors that influence decision-making across the board. Recent research in social networking, virtual communities and recommender systems has focussed on the issue of trust. The majority of recommender system research is carried out on carefully compiled experimental data and does not factor in real world problems such as malicious users. The increasing interest in issues of trust has raised some conflicting opinions on not only its application and relevance, but even its basic definition. Marsh's work in (Marsh 1994) goes some way towards formalising trust in a computational sense, and highlighting different definitions and applications of trust.

Trust in Recommender Systems

There are a number of recent research publications which deal with the issue of how to use trust and reputation models in the recommendation process. This paper focusses on ways to automatically infer trust and reputation from ratings data (using history of errors in contributions to recommendations), and on ways to integrate these models into the recommendation process to produce more reliable recommendations. Other research has focussed on a related issue, but using more invasive methods to acquire trust and reputation models. For example, the work of (Massa & Bhattacharjee 2004) builds a trust model directly from trust data provided by users as part of the popular *epinions.com* service. *Epinions.com* is a web site that allows users to review various items (cars, books, music, etc.). In addition they can assign a trust rating to reviewers based on the degree to which they have found them to be helpful and reliable in the past. (Massa & Bhattacharjee 2004) argue that this trust data can be extracted and used as part of the recommendation process, especially as a means to relieve the sparsity problem that has hampered traditional collaborative

filtering techniques. The sparsity problem refers to the fact that on average two users are unlikely to have rated many of the same items, which means that it will be difficult to calculate their degree of similarity and so limits the range of recommendation partners that can participate in a typical recommendation session. (Massa & Bhattacharjee 2004) argue that it is possible to compare users according to their degree of connectedness in the trust-graph encoded by *Epinions.com*. The basic idea is to measure the distance between two users in terms of the number of arcs connecting the users in the trust-graph encoded by the *Epinions.com* trust data. They show that it is possible to compare far more users according to this method than by conventional forms of ratings similarity and argue that because of this trust-based comparisons facilitate the identification of more comprehensive communities of recommendation partners. However, it must be pointed out that while the research data presented does demonstrate that the trust data makes it possible to compare far more users to each other it has not been shown that this method of comparison maintains recommendation accuracy.

(Montaner, Lopez, & de la Rosa 2002) contemplates the availability of large numbers of virtual recommendation agents as part of a distributed agent-based recommender paradigm. Their main innovation is to consider other agents as personal entities which are more or less reliable or trustworthy and, crucially, that trust values can be computed by pairs of agents on the basis of a conversational exchange in which one agent solicits the opinions of an other with respect to a set of items. Each agent can then infer a trust value based on the similarity between its own opinions and the opinions of the other. Thus this more emphasises a degree of proactiveness in that agents actively seek out others in order to build their trust model which is then used in the *opinion-based* recommendation model. This approach is advantageous from a hybrid recommender perspective, in that agents can represent individual techniques and they can be combined using opinions based on trust.

MODELLING USER TRUST

Recommendation Producers and Consumers

As with previous work in (?) we define two separate profile types in the recommendation process: A *consumer* profile refers to one receiving the item rating, whereas a *producer* refers to the profile that has been selected as a recommendation partner (a contributor to the recommendation) for the consumer and that is participating in the recommendation session. So, to generate a predicted rating for item i for some consumer c , we will typically draw on the services of a number of producer profiles, combining their individual recommendations according to some suitable function, such as Resnick's formula, for example. (see Equation 1). Our benchmark algorithm uses this standard prediction formula (see also (Resnick *et al.* 1994).) In this formula $c(i)$ is the rating to be predicted for item i in consumer profile c and $p(i)$ is the rating for item i by a producer profile p who has rated i , $P(i)$. In addition, \bar{c} and \bar{p} refer to the mean ratings for c and p respectively. The weighting factor $sim(c, p)$ is a measure of the *similarity* between profiles c and p , which is




Figure 1: Calculation of Error Scores from Rating Data Data

traditionally calculated as Pearson's correlation coefficient. One advantage of using this common CF algorithm is that we can compare our approach more easily with existing systems, as seen in our work in (?).

$$c(i) = \bar{c} + \frac{\sum_{p \in P(i)} (p(i) - \bar{p}) \text{sim}(c, p)}{\sum_{p \in P(i)} |\text{sim}(c, p)|} \quad (1)$$

Eliciting Trust From Recommendation Error

We don't believe that the standard method alone is sufficient to produce optimal recommendations. We propose that there is another important element which must be considered in the recommendation process. We refer to this as *trust*. In everyday terms, if a person has successfully recommended lots of items in the past, their recommendations should be awarded more trust than someone who has consistently made poor recommendations in the past.

We define two distinct levels of trust. *Item Level* trust is a score representation for a producers trustworthiness with respect to the recommendation of one specific item. Eg: "John's trustworthiness for recommending Toyota Cars". This is a context-specific trust, as defined in (Abdul-Rahman & Hailes 1997). *Profile Level* trust is a less fine-grained metric, representing a recommendation producers trust as a whole, without respect to one specific item. Eg: "John's trustworthiness". This score is simply an average over the Item Level trust scores for every item in the users profile.

In our work in (?) we calculate trust values at both the Item and Profile levels. Essentially these values summarise the proportion of correct recommendations that a producer has been involved in, according to a pre-defined error bound.

Here, we compute profile and item level error in a similar manner. Consider some producer profile p and some arbitrary item i which has been rated by p . We calculate the average error on p 's contributions to recommendations of item i over a training set of profiles. Each ratings prediction for an item, i , by producer p for a consumer c , is within a finite error ϵ of c 's actual rating $c(i)$; see Equation 4.

In Resnick's approach to CF, producers involved in recommendation would be operating in tandem with a number

of other recommendation partners, so it may not be possible to assess whether or not producer p 's contribution had a positive or negative effect on the final recommendation. To overcome this problem we isolate p 's contribution to the overall recommendation by making p the sole recommendation partner for consumer c for a recommendation on item i . For example, in Figure 1, error scores for item $i1$ are generated for producer b by using the information in profile b only to generate predictions for each consumer profile.

$$T_{p,i,c} = |p(i) - c(i)| = \epsilon \quad (2)$$

Equation 2 shows how each box in Figure 1 contains a separate error score based on each individual profile p being used as the sole recommendation partner for each consumer c and item i .

$$\text{RecSet}(p) = \{(c_1, i_1), \dots, (c_n, i_n)\} \quad (3)$$

The recommendation set for a producer profile p on a specific item i is show in the shaded area of Figure 1, and given by Equation 3. It is worth noting that in a deployed recommender system operating in real time, these values can be generated on-the-fly as new ratings are made in the system. For our experiments, we record 6 different error metrics which we show in the next section can be readily incorporated into the mechanics of a generic CF system to produce a *recommendation range* to be presented to the user.

1. *PError* - The average prediction error over any time p has been involved in producing a recommendation.
2. *IError* - The average prediction error over any time p has been involved in producing a recommendation for each item i .
3. *PUnderest* - The average underestimate over all the times p is involved any recommendation, and there has been an underestimate in the prediction.
4. *IUnderest* - The average underestimate over all the times p is involved in recommendation of each specific item, and there has been an underestimate in the prediction.
5. *POverest* - The average overestimate over all the times p is involved in recommendation of any item, and there has been an overestimate in the prediction.
6. *IOverest* - The average overestimate over all the times p is involved in recommendation of each specific item, and there has been an overestimate in the prediction.

$$IError(p, i) = \frac{\sum_{i \in R} |p(i) - c(i)|}{n} \quad (4)$$

$$PError(p) = \frac{\sum_{i \in P} (IError)}{n} \quad (5)$$

Equation 4 shows our calculation of the basic error for producer profile p at the item level (for item i). n represents the number of recommendations used in the calculation. $p(i)$ is the rating predicted on item i using producer p as the sole recommendation partner for consumer c . $c(i)$ is the actual rating which consumer c gave to item i . This is a fine grained approach to trust value elicitation. A more coarse metric

$PError$ is given by Equation 5, which denotes the average $IError$ over all the items in profile p .

$$IUnderest(p, i) = \frac{\sum_{i \in Rp(i) - c(i)} : if p(i) - c(i) > 0}{n} \quad (6)$$

$$IOverest(p, i) = \frac{\sum_{i \in Rp(i) - c(i)} : if p(i) - c(i) < 0}{n} \quad (7)$$

The case statement in Equation 6 defines the computation of average underestimate at the item level. Here, n represents the number of times there was an underestimate in the predictions generated using producer p for consumer c on item i . A similar computation is denoted in Equation 7 to get the average overestimate in p 's contributions to recommendations of individual items.

To calculate these values at the profile level, we use Equation 8 which is simply an average of the item level underestimates over every item in profile p .

$$PUnderest(p) = \frac{\sum_{i \in P(IUnderest)} : if p(i) - c(i) > 0}{n} \quad (8)$$

We have shown the processes involved in building our trust models based on recommendation error histories. The next step is to demonstrate how these values are manifested in the recommendation process to arrive at a more reliable and transparent recommendation solution.

Using Error Metrics in Recommendation

As a result of these calculations, for every producer profile and every item in these profiles we have the producers normal rating for this item plus the average degree to which this user tends to underestimate ratings, the average degree to which he overestimates ratings, and the overall average error for this producer and item pair. Now how can we use this in recommendation? One obvious direction to take is to produce a *recommendation range*.

For each time we wish to predict a rating for some new target user t and item i we get 3 ratings values: r_1, r_2, r_3 which basically gives us a rating range $[r_2, r_3]$ and an expected rating, r_1 , rather than just a single rating. For example, we can now say to the target user: "My predicted rating for item i is 3.2, but no less than 2.5 and no more than 4" This is showing more information to the target user. Intuitively, this is a better recommendation strategy than simply predicting a set individual value for the user.

We propose three such recommendation strategies:

Addition and Subtraction This is the approach loosely described above. For each recommendation r we generate for an item i , we produce r_1, r_2, r_3 respectively in the form of Equation 9

$$R = \{(r - avg(IUnderest))...r...(r + avg(IOverest))\} \quad (9)$$

The second addition/subtraction approach is more basic and used as our benchmark in comparison with the other techniques. In this approach, instead of using the average

under and overestimates, we simply produce our recommendation range by discounting the standard rating by the average error to get r_1 and increment the standard rating by the average error to arrive at our upper bound r_3 . This is shown in Equation 10

$$R = \{(r - PError)...r...(r + PError)\} \quad (10)$$

In this equation, R is the recommendation range and r is the standard Resnick rating.

Modified Resnick Approach A more interesting way to produce a recommendation range from our error values is to use them *inside* the standard Resnick prediction formula. The basic equation is given in 1, and our modified version in 11, below. Instead of discounting the underestimates after a recommendation has been produced (as with the approach above), we discount the underestimates from all of the neighbours ratings as they are calculated in the standard formula. This produces a lower-bound rating r_u . A similar approach is used to increment neighbours ratings by their pre-calculated average overestimate to produce our upper-bound value r_o .

$$r_u = \bar{c} + \frac{\sum_{p \in P(i)} (p(i) - IUnderest(p, i) - \bar{p}) sim(c, p)}{\sum_{p \in P(i)} |sim(c, p)|} \quad (11)$$

The formula for computing r_o , is that in 11 but with a $+IOverest$ in place of the $-IUnderest$. These results yield another recommendation range:

$$R = \{(r_u)...r...(r_o)\} \quad (12)$$

EVALUATION

In this work we have forwarded a new technique for building a model of user trust based on their history of errors in contribution to recommendations. We have proposed several techniques for integrating these into a standard recommendation algorithm as an additional metric to be considered along with a standard correlation function. In our discussion we will consider the general benefits of producing a recommendation range using these error metrics. The following section describes a specific set of experiments designed to provide us with a better grasp on the empirical improvements our technique can make to the standard CF prediction strategy.

Setup

In this experiment we use the standard MovieLens dataset (Resnick *et al.* 1994) which contains 943 profiles of movie ratings. Profile sizes vary from 18 to 706 with an average size of 105. We divide these profiles into two groups: 80% (753 profiles) are used as the producer profiles and the remaining 20% (190 profiles) are used as the consumer (test) profiles.

Our implementation falls into two phases: Firstly, building the trust-error model, and then using the model in recommendation.

	Avg	Max	Min
IErr	0.938	3.489	0
PErr	0.943	1.804	0.679
IOverest	0.865	3.833	0
IUnderest	0.745	3.489	0
POverest	0.931	2.217	0.328
PUnderest	0.868	1.865	0.162

Table 1: Error Information

Building the Error-Trust Model In a deployed recommender system, error values can be computed in real time, as users enter ratings to the system. In this experiment however, we build the trust values separately. For each producer profile in our training set, we temporarily assign that profile to position of *consumer*, and using the standard Resnick prediction formula, we generate predictions for that consumer, using each of the other producer profiles in turn as the sole recommendation partner. We assess the proximity of the predicted ratings to the actual ratings by employing a standard leave-one-out test, where we hide $n\%$ of the consumers profile and try to predict the hidden ratings. We compare the predicted ratings to the known real rating and record the error. These errors fall into either underestimate or overestimate categories.

This procedure is carried out for every producer-item pair in the system, and the result is an average error for each producer-item pair over all of the times that producer was involved in the production of a recommendation of that item. This is our item-level model. In order to get our more coarse grained profile-level model, we average these error values again, but this time over each item in a producers profile. This provides us with an average error per producer (profile level error).

Table gives us an overview of the type of recommendation errors we found for each of our strategies. Noting that IErr and PErr values are for overall average error at the item and profile levels respectively, and do not necessarily have to be the sum of the mean individual under and overestimates at that level. The largest average error is at the item level at 0.937 on the Movielens rating scale of 1 to 5. This is as expected from our recommendation results in the next section, in which the recommendation ranges produced based on these errors outperforms the others in accuracy tests, most likely due to the broader range produced from these values. We note that the profile level errors are much less extreme, *POverest*, for example having a minimum error of 0.328 in comparison with minimum *IOverest* at 0 error. This result is exactly as expected since our profile level errors are averaged over every item in a users profile. A similar trend is the case for the other profile / item level comparisons.

Figure 2 depicts the distribution of the error values at item and profile levels respectively. Both follow a normal distribution. Profile level errors are centered more closely around the mean of 0.94, having a standard deviation of 0.13, compared with 0.3 around a mean of 0.94 for the item level er-




Figure 2: The distribution of producer error values

rors. This variance should have a notable effect when they are used to generate our recommendation ranges in the next section, yielding potentially more broadness for the item level approach. Space restrictions prevent us from showing histograms for each individual technique described, but the others also follow normal distributions, with a similar variance between profile and item levels.

Predictive Accuracy Experiment.

The overarching goal of this research is to improve predictive accuracy for standard CF based recommendation algorithms. This experiment looks at the accuracy of our technique from a % correct perspective. For the 190 profiles in our test set, we generated recommendation ranges for each rated item using the three techniques described in section 3, using both item level and profile level error values separately for our modified resnick approach, and an average error value for the addition / subtraction approach. For this set we defined accuracy as a consumers actual rating falling within the predicted recommendation range. Note that training and test data are completely independent sets. Results of this experiment are presented in Figure 4. It is clear that the modified resnick approach using item level error values outperforms the other techniques, beating its closest rival *avgErr* by 4.5%. This may be due to the greater variance in the item level errors. The worst performer in fact is the modified resnick approach using profile level errors, at 14% worse performance than its item level counterpart. This suggests that the variance in errors does have a notable effect on predictive accuracy for collaborative filtering.

Comparison with a Benchmark Resnick Algorithm

Due to the *recommendation range* nature of our predictions, we must define a new method to evaluate the performance of our technique with respect to a benchmark resnick algorithm. To this end, we define a scalar constant which is the average of the absolute differences between our upper bound r_o and our lower bound r_u , (which is 1.85). We propose that a fair comparison would be to assume a standard resnick prediction to be 'correct' if it falls within half of this distance either side of the real rating. This is only a pseudo-comparison, as it is impossible to compare these techniques directly. From Figure 4, it is clear that the range technique performs 15% better than the benchmark using this comparison. This does show that the direction in which the error models pull the r_o and r_u ratings is having a positive effect




Figure 3: Distribution of the recommendation ranges and actual ratings for the modified Resnick technique




Figure 4: % Ratings inside recommendation range

on predictive accuracy.

Figure 3 is a trend graph of the of the recommendation ranges for our best performer, the modified Resnick approach using item level errors. Here we can see the upper bounds r_o and lower bounds r_u of the recommendation range generated by using $IUnderest$ and $IOverest$ respectively in the standard Resnick prediction formula, see Equations 1 and 11. The real ratings provided by our test profiles tend to remain within the range, even when the standard Resnick prediction tends away from the real rating. This outcome has a positive effect not only on predictive accuracy, but also on users overall trust in the recommender system as a whole, reducing the mal-effects of false positives in recommendations.

DISCUSSION & CONCLUSIONS

Trust & CF Robustness The work of O'Mahony et. al (O'Mahony *et al.*) and Leiven (Levien) outlines recommender systems from the viewpoints of *accuracy*, *efficiency*, and *stability*. The trust-error models defined in this paper can not only be used to increase recommendation accuracy, they can be utilised to increase the overall robustness of CF systems. (O'Mahony *et al.*) defines several attack strategies that can adversely skew the recommendations generated by a K-NN CF system. They show empirically in (O'Mahony *et al.*) that a CF system needs to attend to each of these factors in order to succeed well. There are many motivations for users attempting to mislead recommender systems, including profit and malice, as outlined in (Lam & Riedl 2004). The trust-error metrics presented in this paper can

help a system detect malicious users, as they will most likely tend to have higher errors in their contribution histories as producers. Using a weighting scheme we describe in (?) we can render their contributions to recommendation ineffective based on their reputation values. work by Sinha and Swearingen shows the importance of transparency in recommender system interfaces, (Sinha & Swearingen 2002). Our recommendation range prediction method shows more information to the user about what the system knows, thereby increasing user trust in the recommender. In a similar vein, we can also mention the trustworthiness of the contributing producers to the consumer.

To conclude, here we have proposed two approaches to building up trust values based on error models, one that operates with respect to a specific user-item combination, eg: 'John's reputation for recommendation Toyota Carinas', and another that operates at a user level; eg: 'John's trustworthiness as a recommender'. We presented three techniques for integrating these models into a standard CF algorithm to produce recommendation *ranges* to present to a user, and defined a fair method for comparing our approach to the benchmark. Our empirical results indicate that trust-aided recommendation based on error models increases predictive accuracy over the benchmark by 15%.

ACKNOWLEDGMENTS

This material is based on works supported by
***** under Grant No. *****

References

- Abdul-Rahman, A., and Hailes, S. 1997. A distributed trust model. In *New Security Paradigms 1997*, 48–60.
- Burke, R. D.; Hammond, K. J.; and Young, B. C. 1997. The findme approach to assisted browsing. *IEEE Expert* 12(4):32–40.
- Dellarocas, C. Building trust on-line: The design of reliable reputation reporting mechanisms for online trading communities. *eBusiness@MIT*, July 2001.
- Lam, S. K., and Riedl, J. 2004. Shilling recommender systems for fun and profit. In *Proceedings of the 13th international conference on World Wide Web*, 393–402. ACM Press.
- Levien, R. Attack resistant trust metrics. *Ph.D Thesis, UC Berkeley*.
- Marsh, S. 1994. Formalising trust as a computational concept. *Ph.D. Thesis. Department of Mathematics and Computer Science, University of Stirling*.
- Massa, P., and Bhattacharjee, B. 2004. Using trust in recommender systems: an experimental analysis. *Proceedings of 2nd International Conference on Trust Management, Oxford, England*.
- Montaner, M.; Lopez, B.; and de la Rosa, J. L. 2002. Developing trust in recommender agents. In *Proceedings of the first international joint conference on Autonomous agents and multiagent systems*, 304–305. ACM Press.

O'Mahony, M.; Hurley, N.; Kushmerick, N.; and Silvestre, G. Collaborative recommendation: A robustness analysis, url: citeseer.ist.psu.edu/508439.html.

Resnick, P.; Iacovou, N.; Suchak, M.; Bergstrom, P.; and Riedl, J. 1994. Grouplens: An open architecture for collaborative filtering of netnews. In *Proceedings of ACM CSCW'94 Conference on Computer-Supported Cooperative Work, Sharing Information and Creating Meaning*, 175–186.

Sinha, R., and Swearingen, K. 2002. The role of transparency in recommender systems. In *CHI '02 extended abstracts on Human factors in computing systems*, 830–831. ACM Press.